

Homework 3!

March 7, 2016

1 Welcome to homework 3 by Willy G

Let's do some imports.

```
In [1]: %matplotlib inline
```

```
import scipy.io
from sklearn import svm
from sklearn.preprocessing import normalize
import numpy as np
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
from scipy.stats import multivariate_normal
```

```
from skimage.feature import hog
from skimage import data, color, exposure
```

```
In [26]: delta = 0.25
x = np.arange(-5.0, 5.0, delta)
y = np.arange(-5.0, 5.0, delta)
X, Y = np.meshgrid(x, y)

prob2mu = [np.array([1,1]),
            np.array([-1,2]),
            np.array([0,2]),
            np.array([0,2]),
            np.array([0,2])]

prob2cov = [np.array([[2, 0],
                      [0, 1]]),
            np.array([[3, 1],
                      [1, 2]]),
            np.array([[1, 1],
                      [1, 2]]),
            np.array([[2, 1],
                      [1, 2]]),
            np.array([[1, 1],
                      [1, 2]])]


for mu,cov in zip(prob2mu, prob2cov):
    plt.rcParams['xtick.direction'] = 'out'
    plt.rcParams['ytick.direction'] = 'out'
```

```

gaussian = multivariate_normal(np.array(mu), np.array(cov), True)
z = [np.array([p1[idx],p2[idx], gaussian.pdf(np.array([p1[idx],p2[idx]])))] for idx, p1 in enumerate(p1s)]

plt.contour(np.transpose(z))          # line colours

```



Time to get some awesome data for MNIST and do some dundiddly preprocessing.

```

In [12]: def featurize(image, i ):
    lol = image.reshape(28,28)
    extra1 = hog(lol, orientations=8, pixels_per_cell=(14, 14), cells_per_block=(1, 1), visualis
    extra2 = hog(lol, orientations=8, pixels_per_cell=(7, 7), cells_per_block=(1, 1), visualis
    pixelvector = list(image* 1/255.0*2 - 1)
    if (i-1) % 10000 == 0:
        print("Featurization progress: %s" % i)
    pixelvector.extend(list(extra1[1].flatten()*1/255.0*4))
    pixelvector.extend(list(extra2[1].flatten()*1/255.0*4))
    return np.array(pixelvector)

print("Starting!")

image_train = scipy.io.loadmat("../data/digit_dataset/train.mat")
image_test = scipy.io.loadmat("../data/digit_dataset/test.mat")['test_images']

image_train_data_raw = image_train['train_images']
image_train_label_raw = image_train['train_labels']

print("Data loaded!")
image_tdata = np.array([featurize(image_train_data_raw[:, :, i].flatten(), i) for i in range(len
image_tlabel = image_train_label_raw.ravel()

```

```

image_test_data = np.array([featurize(test.reshape(28,28).T.flatten(), i) for i, test in enumerate(test_data)])

#Shuffle that image data good.
shuffle = np.random.permutation(np.arange(image_tdata.shape[0]))
image_tdata, image_tlabel = image_tdata[shuffle], image_tlabel[shuffle]

#VALIDATION
image_valid_data = image_tdata[0:10000]
image_valid_label = image_tlabel[0:10000]

#TRAINING
image_train_data =image_tdata[10000:]
image_train_label =image_tlabel[10000:]

plt.matshow(image_train_data_raw[:, :, 1])
print(image_train_label[0]) #unrelated to the plot!

```

Starting!

Data loaded!

Featurization progress: 1

Featurization progress: 10001

Featurization progress: 20001

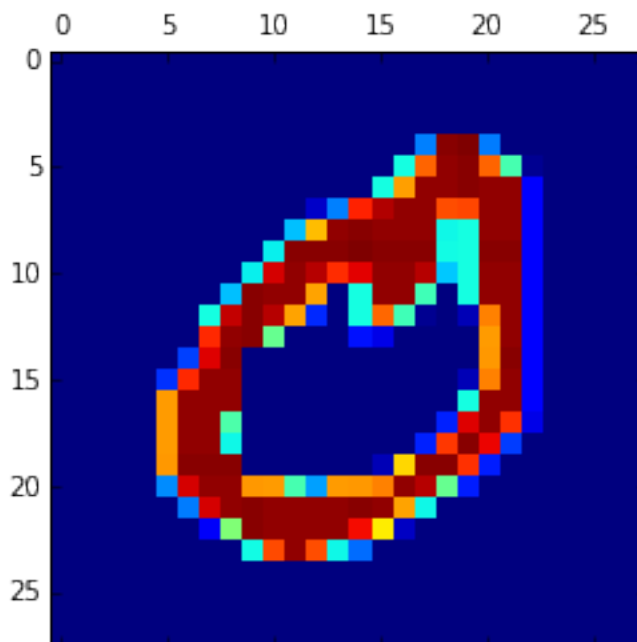
Featurization progress: 30001

Featurization progress: 40001

Featurization progress: 50001

Featurization progress: 1

6



Did you like that. I did. Okay so, let's make a class which model data classes!

```

In [13]: def perturb_eig(mat, alpha):
          return mat + alpha*np.eye(mat.shape[0])

class dataclass:
    def __init__(self, datum, labels, class_label, alpha):
        # Find all of the datapoints in datum with the labels class_lab.
        # add them to this dataclass
        ids_classc = [idx for idx,label in enumerate(labels) if label == class_label ]
        self.data = datum[ids_classc]

        #calculate some things!
        singcov = np.cov(self.data.T)
        self.mean = np.mean(self.data, axis = 0)
        self.prior = len(ids_classc)/len(labels)

        #Let's make sure we can invert the covariance matrix using a bincary searach.
        self.cov = singcov

        self.gaussian = multivariate_normal(self.mean, self.cov, True)
        self.covall = None

    def lda(self, covall, x):
        if self.covall is covall:
            return self.ldagaussian.logpdf(x) + self.prior
        else:
            self.covall =covall
            self.ldagaussian = multivariate_normal(self.mean, covall, True)
            return self.ldagaussian.logpdf(x) + self.prior

    def qda(self, x):
        return self.gaussian.pdf(x)*self.prior

```

We gotta make sure this class makes sense

```



In [17]: a = dataclass(image_train_data, image_train_label, 3.0, 0.1)


plt.matshow(a.data[0][0:28*28].reshape(28,28))
plt.matshow(a.data[0][28*28*2:].reshape(28,28))
plt.matshow(a.cov)


plt.matshow(a.cov[1333,:].reshape(56+28,28))

Out[17]: <matplotlib.image.AxesImage at 0x7f72d5d8a4a8>

```








Let's actually make all of the image classes!

```
In [18]: image_train_class = [dataclass(image_train_data, image_train_label, float(x), 0.1) for x in range(10000)]
```

```
In [19]: plt.matshow(image_train_class[7].mean.reshape((56+28,28)))
```


```
Out[19]: <matplotlib.image.AxesImage at 0x7f72d88f9048>
```





Well that worked thank god! Now we need to add all of those peski general properties

```
In [20]: covall = 0.1*np.sum([image_train_class[i].cov for i in range(len(image_train_class))], axis=0)
         precall = np.linalg.pinv(covall)
         plt.matshow(covall)
         plt.matshow(precall)
```

Out[20]: <matplotlib.image.AxesImage at 0x7f72d8c2f5f8>





2 Question 5 part(d)

Now we make some helper methods.

```
In [21]: def best_prediction(classes,method, x):
    return np.argmax(list(map(lambda cls: method(cls, x), classes)))

def validate(classes, method, x, y):
    if int(best_prediction(classes, method, x)) == int(y):
        return 1.0
    else:
        return 0

def lda(cls, x):
    return cls.lda(covall, x)

def qda(cls, x):
    return cls.qda(x)

def experiment(method, num_class, train_data, train_labels, validation, labels, samples=100000):
    classes = [dataclass(train_data[0:samples], train_labels[0:samples], float(x), 0.1) for x in range(num_class)]
    print("classes created for experiment!")
    net = 0
    for idx, sample in enumerate(validation):
        score = validate(classes, method, sample, labels[idx])
        net += score

    return 1.0 - net/len(validation)
```

```

def do_test(method, num_class, train_data, train_labels, test_data):
    classes = [dataclass(train_data, train_labels, float(x), 0.1) for x in range(num_class)]
    print("classes classes created for test!")
    pred = [int(best_prediction(classes, method, test)) for test in test_data]
    pairs = list(enumerate(pred))
    for idx, pair in enumerate(pairs):
        pairs[idx] = (pair[0]+1, pair[1])
    return pairs

experiment_count = [100,200, 500, 1000, 2000, 5000, 10000, 30000, 50000]

```

Now we can actually run some experiments

```

In [22]: #LDA
print("LDA")
errors = []
for count in experiment_count:
    print("Running at %s," % count)
    error = experiment(lda, 10, image_train_data,image_train_label , image_valid_data, image_v
    print("\tError: %s" % error)
    errors.append(error)

plt.plot(experiment_count, errors)
plt.show()


```

LDA

```

Running at 100,
classes created for experiment!
Error: 0.6674
Running at 200,
classes created for experiment!
Error: 0.36329999999999996
Running at 500,
classes created for experiment!
Error: 0.13729999999999998
Running at 1000,
classes created for experiment!
Error: 0.08689999999999998
Running at 2000,
classes created for experiment!
Error: 0.06969999999999998
Running at 5000,
classes created for experiment!
Error: 0.0610000000000000054
Running at 10000,
classes created for experiment!
Error: 0.057400000000000001
Running at 30000,
classes created for experiment!
Error: 0.0542000000000000026
Running at 50000,
classes created for experiment!
Error: 0.054100000000000004

```



```
In [23]: #QDA
print("QDA")
errors = []
for count in experiment_count:
    print("Running at %s," % count)
    error = experiment(qda, 10, image_train_data, image_train_label, image_valid_data, image_v
    print("\tError: %s" % error)
    errors.append(error)

plt.plot(experiment_count, errors)
plt.show()
```

```
QDA
Running at 100,
classes created for experiment!
Error: 0.9071
Running at 200,
classes created for experiment!
Error: 0.8905
Running at 500,
classes created for experiment!
Error: 0.7746
Running at 1000,
classes created for experiment!
Error: 0.2619
Running at 2000,
classes created for experiment!
Error: 0.17300000000000004
Running at 5000,
```

```


classes created for experiment!
Error: 0.7823
Running at 10000,
classes created for experiment!
Error: 0.2914
Running at 30000,
classes created for experiment!
Error: 0.14239999999999997
Running at 50000,
classes created for experiment!
Error: 0.13470000000000004

```

```

/home/william/anaconda3/lib/python3.4/site-packages/scipy/stats/_multivariate.py:603: RuntimeWarning: ov
return np.exp(self.logpdf(x))

```



```

In [27]: thedata = list(do_test(lda, 10, image_tdata, image_tlabel, image_test_data))
np.savetxt(
    'kaggleimage.csv',          # file name
    thedata,                    # array to save
    fmt='%i',                   # formatting, 2 digits in this case
    delimiter=',',             # column delimiter
    newline='\n')

```

classes created for test!

Now let's do the spam thing and work it all out!

```

In [25]: spam_data = scipy.io.loadmat("../data/spam_dataset/spam_data.mat")
spam_train = spam_data['training_data']
spam_test = spam_data['test_data']

```

```

spam_train_data_raw = spam_train
spam_train_label_raw = spam_data['training_labels']

spam_tlabel = spam_train_label_raw.ravel()
div_train = [1.0/max(arr) for arr in spam_train_data_raw.T]
spam_tdata = np.array([np.multiply(dp, div_train) for dp in spam_train_data_raw ])
spam_test_data = np.array([np.multiply(dp, div_train) for dp in spam_test ])

#Shuffle that spam data good.
shuffle = np.random.permutation(np.arange(spam_tdata.shape[0]))
spam_tdata, spam_tlabel = spam_tdata[shuffle], spam_tlabel[shuffle]

#VALIDATION
spam_valid_data = spam_tdata[0:750]
spam_valid_label = spam_tlabel[0:750]

#TRAINING
spam_train_data =spam_tdata[750:]
spam_train_label =spam_tlabel[750:]

```

Looks good! Let's train on everything and build a Kaggle.

```

In [26]: thedata = list(do_test(qda, 2, spam_tdata, spam_tlabel, spam_test_data))
np.savetxt(
    'kagglespam.csv',          # file name
    thedata,                   # array to save
    fmt='%i',                  # formatting, 2 digits in this case
    delimiter=',',            # column delimiter
    newline='\n')

```

classes created for test!

3 Linear Regressive

Let's do it! We know that $\langle w, x \rangle = \hat{y}$. Therefore, we can consider the following line of reasoning.

Let X be the design matrix containing all of the data points. Since for each data point there is weight whose dot product yields an entry of the final output y , we let w be some weight vector, along which we will multiply each data point. This gives

$$Xw = y.$$

Performing a pseudo inverse will give us the least squares solution $w = (X^T X)^{-1} X^T y$.

```

In [142]: def linreg(data, y, alpha=0):
    #Add one to all the data points.
    extended = np.array([np.append(sample, [1]) for sample in data])
    print(extended)
    w = np.dot(np.dot(np.linalg.inv(np.dot(extended.T, extended)), extended.T), y)
    return w

def linrss(data, y, w):
    extended = np.array([np.append(sample, [1]) for sample in data])
    return np.linalg.norm(np.dot(extended,w) - y,2)**2

def plot_weights(w):
    plt.bar(range(len(w)-1),w[:-1])

```

```

In [143]: test = np.array([[0],[1],[1.2],[1.5],[2]])
          test_y = np.array([1.8,1.6,1.5,1.3,1.2])

          w = linreg(test,test_y, 0)
          #plot_weights(w)

          X = np.array([[-1,1],
                        [-0.5,1],
                        [0,1],
                        [0.5,1],
                        [1,1],
                        [1.5,1],
                        [2,1],
                        [2.5,1],
                        [3,1]])
          x = np.array([-1,-0.5,0,0.5,1,1.5,2,2.5,3])

          plt.plot(x, np.dot(X,w))
          plt.plot(test, test_y, 'ro')


          print(linrss(test, test_y, w))

```

```

[[ 0.  1. ]
 [ 1.  1. ]
 [ 1.2  1. ]
 [ 1.5  1. ]
 [ 2.  1. ]]
0.0133120437956

```



```

In [149]: real_data = scipy.io.loadmat("../data/housing_dataset/housing_data.mat")
          real_Xtrain = (real_data['Xtrain'])/ np.linalg.norm((real_data['Xtrain']),2)

```



```

real_Ytrain = real_data['Ytrain'].ravel()/np.linalg.norm(real_data['Ytrain'].ravel(),2)
real_Xvalidate = real_data['Xvalidate'] /np.linalg.norm((real_data['Xvalidate'])),2)
real_Yvalidate = real_data['Yvalidate'].ravel()/np.linalg.norm(real_data['Yvalidate'].ravel()

```

In [151]: `w = linreg(real_Xtrain, real_Ytrain)`

```


plot_weights(w)
print(linrss(real_Xvalidate, real_Yvalidate, w))

```

```

[[ 5.32672469e-06  5.82932150e-05  1.97286100e-03 ...,  6.92960594e-05
 -2.20621602e-04  1.00000000e+00]
 [ 8.92487339e-06  6.19365410e-05  6.02241778e-03 ...,  6.92049762e-05
 -2.23244797e-04  1.00000000e+00]
 [ 9.20449865e-06  4.73632372e-05  4.10967166e-03 ...,  6.94600090e-05
 -2.20967718e-04  1.00000000e+00]
 ...,
 [ 2.73249445e-06  7.10448558e-05  1.84898791e-03 ...,  6.79480288e-05
 -2.19473955e-04  1.00000000e+00]
 [ 1.69068540e-05  5.64715521e-05  4.94945829e-03 ...,  6.80573285e-05
 -2.22388615e-04  1.00000000e+00]
 [ 5.94208244e-06  9.47264744e-05  6.59077662e-03 ...,  6.89499434e-05
 -2.22716515e-04  1.00000000e+00]]
613157.990429

```




In [154]: `extended = np.array([np.append(sample, [1]) for sample in real_Xvalidate])`
`plt.hist(np.dot(extended, w) - real_Yvalidate)`

```

Out[154]: (array([ 4., 21., 102., 423., 553., 81., 6., 4., 5., 1.]),
 array([ 20.43552208, 20.97386487, 21.51220766, 22.05055045,
        22.58889324, 23.12723604, 23.66557883, 24.20392162,

```

```
24.74226441, 25.28060721, 25.81895  ]),  
<a list of 10 Patch objects>)
```



In []: