

CS 189: Homework 4

William Guss
26793499
wguss@berkeley.edu

April 1, 2016

1. Ridge Regression

- (a) Let $J(w, \alpha)$ be a loss function so that

$$J(w, \alpha) = (Xw + \alpha 1 - y)^T (Xw + \alpha 1 - y) + \lambda w^T w. \quad (1)$$

The maximizing w, α are derived as follows. The gradient with respect to w is calculated; assuming that J is convex in w , this will give a maximum. A proof is given at the end of the assignment.

$$\begin{aligned} \nabla_w J &= \nabla_w (Xw + \alpha 1 - y)^T (Xw + \alpha 1 - y) + \nabla_w \lambda w^T w \\ &= 2\lambda w + 2(\nabla_w (Xw + \alpha 1 - y))^T (Xw + \alpha 1 - y) \\ &= 2\lambda w + 2X^T (Xw + \alpha 1 - y) \end{aligned}$$

Letting this quantity be 0 by the convexity of J we get

$$\begin{aligned} 0 &= 2\lambda w + 2X^T (Xw + \alpha 1 - y) \\ 0 &= 2\lambda w + 2X^T Xw + 2\alpha X^T 1 - 2X^T y \\ 0 &= (\lambda I + X^T X)w + \alpha X^T 1 - X^T y \\ X^T y - \alpha \bar{x} &= (\lambda I + X^T X) \\ w &= (\lambda I + X^T X)^{-1} X^T y \end{aligned} \quad (2)$$

- (b) See attached iPython notebook!

2. Logistic Regression

- (a) In this case

$$R(w) = 1.9883724141284103. \quad (3)$$

- (b) Nah

(c) $w(1) = \text{array}([-3., 5.05089812, 0.68363271])$

(d) $R(w(1)) = 0.066133848540594925$

- (e) Nah

(f) $w(2) = \text{array}([-4., 9.10179623, 1.36726541])$

(g) $R(w(2)) = 0.0015778803918752944$

3.

4. Revisiting Logistic Regression

(a) Show the following theorem:

Theorem 1. *If we define $g(z) = \frac{\tanh z + 1}{2}$ then*

$$g(z) = \frac{e^z - e^{-z}}{2(e^z + e^{-z})} + \frac{1}{2}. \quad (4)$$

Proof. Consider the following algebraic manipulation of $g(z)$ for any $z \in \mathbb{R}$.

$$\begin{aligned} g(z) &= \frac{\tanh z + 1}{2} = \frac{(2s(2z) - 1) + 1}{2} \\ &= \frac{2s(2z) - 1}{2} + \frac{1}{2}. \end{aligned}$$

Therefore we must only show that

$$\frac{e^z - e^{-z}}{2(e^z + e^{-z})} = \frac{2\frac{1}{1+e^{-2z}} - 1}{2}. \quad (5)$$

Observe that since $e^z \neq 0$,

$$\begin{aligned} \frac{2s(2z) - 1}{2} &= \frac{1 - e^{-2z}}{2(1 + e^{-2z})} \\ &= \frac{e^z - e^{-z}}{2(e^z + e^{-z})} \end{aligned}$$

and so we know that,

$$g(z) = \frac{e^z - e^{-z}}{2(e^z + e^{-z})} + \frac{1}{2}. \quad (6)$$

□

(b) We calculate $g'(z)$ as follows,

$$g'(z) = \frac{1}{2} \tanh' z = \frac{1}{2} (1 - \tanh^2 z)$$

by the definition of $s'(z)$.

(c) Let $J(w) = \sum_{i=1}^n y_i \ln(g(X_i \cdot w)) + (1 - y_i) \ln(1 - g(X_i \cdot w))$. We derive the batch gradient descent learning rule as follows.

$$\begin{aligned}
\nabla J(w) &= \sum_{i=1}^n y_i \nabla \ln(g(X_i \cdot w)) + \nabla (1 - y_i) \ln(1 - g(X_i \cdot w)) \\
&= \sum_{i=1}^n \frac{y_i}{g(X_i \cdot w)} \nabla g(X_i \cdot w) - \frac{1 - y_i}{1 - g(X_i \cdot w)} \nabla g(X_i \cdot w) \\
&= \sum_{i=1}^n \frac{y_i}{g(X_i \cdot w)} g(X_i \cdot w) (1 - g(X_i \cdot w)) \nabla X_i \cdot w - \\
&\quad \frac{1 - y_i}{1 - g(X_i \cdot w)} g(X_i \cdot w) (1 - g(X_i \cdot w)) \nabla X_i \cdot w \\
&= \sum_{i=1}^n \left(\frac{y_i}{g(X_i \cdot w)} - \frac{1 - y_i}{1 - g(X_i \cdot w)} \right) g(X_i \cdot w) (1 - g(X_i \cdot w)) \nabla X_i \cdot w \\
&= \sum_{i=1}^n \left(\frac{y_i}{g(X_i \cdot w)} - \frac{1 - y_i}{1 - g(X_i \cdot w)} \right) g(X_i \cdot w) (1 - g(X_i \cdot w)) X_i \\
&= \sum_{i=1}^n \left(\frac{y_i(1 - g(X_i \cdot w)) - (1 - y_i)g(X_i \cdot w)}{g(X_i \cdot w)(1 - g(X_i \cdot w))} \right) g(X_i \cdot w) (1 - g(X_i \cdot w)) X_i \\
&= \sum_{i=1}^n (y_i(1 - g(X_i \cdot w)) - (1 - y_i)g(X_i \cdot w)) X_i \\
&= \sum_{i=1}^n (y_i - y_i g(X_i \cdot w) - g(X_i \cdot w) + y_i g(X_i \cdot w)) X_i \\
&= \sum_{i=1}^n (y_i - g(X_i \cdot w)) X_i
\end{aligned}$$

Therefore for every epoch perform the following weight update rule.

$$w \leftarrow w - \lambda \sum_{i=1}^n (y_i - g(X_i \cdot w)) X_i. \quad (7)$$

It is important to note that this loss function although rooted in a probabilistic model, has some other quantitatively advantageous properties. For example, if square loss function were used, the Δw term would be limited, the gradient would vanish as $w \cdot X_i$ approached large negative or positive values, which in some cases may not be advantageous. This avoids such vanishing gradient problems in a manner that again is rooted in the probabilistic model.

5. **Daniel's Email Problem.** For this problem we will denote $X \subset \mathbb{R}^n$ as the manifold to which the feature vectors for all probable email scenarios are close. Let x_t denote the component of some $x \in X$ such that x_t describes Daniel's time since midnight feature.

From the problem it is observed that the restriction of X to the dimension of x_t is essentially dense in $S = [0, \epsilon_1] \sqcup (\epsilon_2, 24]$ for x in the spam class. Clearly in the dimension of x_t , this is not a linearly separable problem. One might consider a

quadratic kernel allowing the classification of S^C inside of the maximal boundary of S with respect to the origin centered norm ordering of \mathbb{R} , (take for example (ϵ_1, ϵ_2)). However a simpler solution arises by changing the feature itself.

Let $d : \mathbb{R} \times \mathbb{R} \rightarrow \mathbb{R}$ be the normal euclidean distance on \mathbb{R} . Then defining the feature x_t so that

$$x_t = \min\{d(x, 0), d(x, 24)\}. \quad (8)$$

This function then implies that the restriction of X to the dimension of x_t has a roughly contiguous interval of spam (positive) samples, ie. $[0, \max\{\epsilon_1, \epsilon_2\}]$. This interval is of course linearly separable with respect to $[0, 24]$, and therefore we can use a Linear SVM.

Homework 4

April 1, 2016

Problem 4

```
In [33]: import numpy as np
import math

In [55]: def s(gamma):
    return 1.0/(1.0 + math.exp(-gamma))

    def R(w, X, y):
        net = 0
        for i in range(X.shape[0]):
            print(X[i])
            net += y[i] * math.log(s(np.dot(w, X[i]))) + (1-y[i])*math.log(s(np.dot(w, X[i])))
        return -net

In [56]: X = np.array([[0,3,1],
                        [1,3,1],
                        [0,1,1],
                        [1,1,1]])
y = np.array([1,1,0,0])
w0 = np.array([-2,1,0])

In [57]: R(w0,X,y)

[0 3 1]
[1 3 1]
[0 1 1]
[1 1 1]

Out[57]: 1.9883724141284103

In [58]: def delta_w(w,X,y):
    net = np.zeros(w.shape[0])
    for i in range(X.shape[0]):
        net += y[i] - s(np.dot(X[i], w))*X[i]
    return -net

In [59]: w1 = w0 + delta_w(w0, X,y)
w1

Out[59]: array([-3.          ,  5.05089812,  0.68363271])

In [60]: R(w1,X,y)

[0 3 1]
[1 3 1]
[0 1 1]
[1 1 1]
```

```
Out[60]: 0.066133848540594925
```

```
In [61]: w2 = w1 + delta_w(w0, X,y)
          w2
```

```
Out[61]: array([-4.          ,  9.10179623,  1.36726541])
```

```
In [62]: R(w2,X,y)
```

```
[0 3 1]
```

```
[1 3 1]
```

```
[0 1 1]
```

```
[1 1 1]
```

```
Out[62]: 0.0015778803918752944
```

```
In [ ]:
```