

CS 189: Homework 7

William Guss
26793499
wguss@berkeley.edu

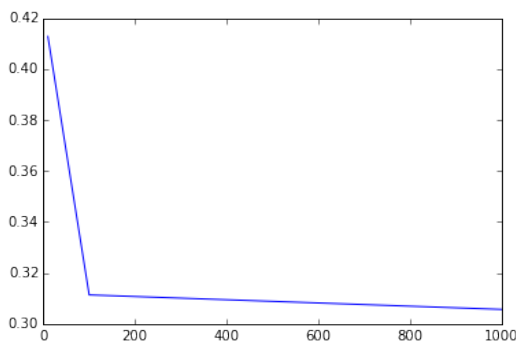
May 3, 2016

1. KMeans

- (a) After running K -Means Figure 1 is a sample plot of each class. Furthermore running K -means with different initializations leads to different loss; the classes are distributed differently each time.

2. Joke Recommender System

- (a) *Warm Up*. We tried the average value recommendation system with some success! We got around 37% error as seen in the iPython notebook.
- (b) k -NN. We tried using k -NNS and below is a plot of the errors we received. k -NNs did best at $k = 1000$. Awesome!



- (c) *Latent Factor Model*. We apply the latent factor model by setting all the values in the matrix to 0 where there are NaN s. We then perform SVD yielding

$$R = USV^T, \quad u_i = U_i, \quad v_j = (SV^T)_j^T \quad (1)$$

Then we let $R_{ij} = \langle u_i, v_j \rangle$. Below is a plot of how well the MSE of this approximation does with $S = \text{diag}(s_1, \dots, s_d, 0 \dots)$ as $d \rightarrow 100$.


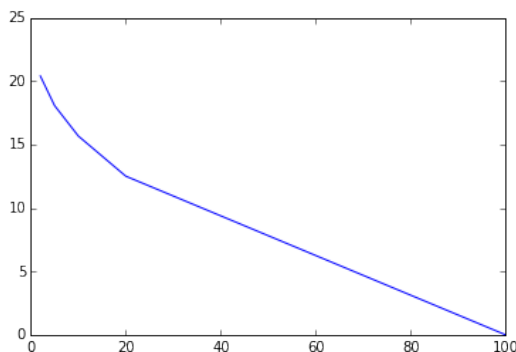
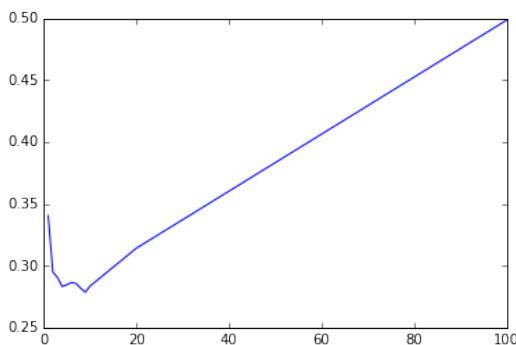


Figure 1: The mean centers of each class.



Applying these approximations to the validation set we get the following graph indicating that 10 latent dimensions optimizes the proper embedding of the data.



- (d) We will now attempt to do L_1 regularized gradient descent on the Frobenius MSE. Notice that $L(u_i, v_j)$ is symmetric in its calculation of the gradient so we skip the second derivations. First we calculate the gradient with respect to a single u_i .

$$\begin{aligned}
 \frac{\partial L}{\partial u_i} &= \frac{\partial}{\partial u_i} \sum_{(k,l) \in S} (v_l^T u_k - R_{kl})^2 + \frac{\partial}{\partial u_i} \lambda \sum_{k=1}^d u_k^T u_l \\
 &= \frac{\partial}{\partial u_i} \sum_{k \in S_1} \sum_{l \in S_2} (v_l^T u_k - R_{kl})^2 + 2\lambda u_i \\
 &= \frac{\partial}{\partial u_i} \sum_{l \in S_2} (v_l^T u_i - R_{il})^2 + 2\lambda u_i \\
 &= 2 \sum_{l \in S_2} (v_l^T u_i - R_{il}) v_l + 2\lambda u_i
 \end{aligned}$$

Using the aforementioned symmetry we get

$$\begin{aligned}
 \frac{\partial L}{\partial u_i} &= \sum_{l \in S_2} (v_l^T u_i - R_{il}) v_l + \lambda u_i \\
 \frac{\partial L}{\partial v_j} &= \sum_{k \in S_1} (v_j^T u_k - R_{kj}) u_k + \lambda v_j
 \end{aligned} \tag{2}$$

Now we calculate the minimizers

$$\begin{aligned}
 \frac{\partial L}{\partial u_i} = 0 &= \sum_{l \in S_2} (v_l^T u_i - R_{il}) v_l + \lambda u_i \\
 0 &= \lambda I u_i + \sum_{l \in S_2} v_l v_l^T u_i - \sum_{l \in S_2} v_l R_{il} \\
 \sum_{l \in S_2} v_l R_{il} &= \left(\lambda I + \sum_{l \in S_2} v_l v_l^T \right) u_i \\
 u_i &= \left(\lambda I + \sum_{l \in S_2} v_l \otimes v_l \right)^{-1} \sum_{l \in S_2} v_l R_{il}
 \end{aligned}$$

This gives the following alternating minimizer method for every i, j alternate the following.

$$\begin{aligned}
 u_i &= \left(\lambda I + \sum_{l \in S_2} v_l \otimes v_l \right)^{-1} \sum_{l \in S_2} v_l R_{il} \\
 v_j &= \left(\lambda I + \sum_{k \in S_1} u_k \otimes u_k \right)^{-1} \sum_{k \in S_1} u_k R_{kj}
 \end{aligned} \tag{3}$$

We implemented this in iPython and got the following descent for $\lambda = 10$.

